

# Character\_Level\_RNN\_Exercise

May 15, 2020

## 1 Character-Level LSTM in PyTorch

In this notebook, I'll construct a character-level LSTM with PyTorch. The network will train character by character on some text, then generate new text character by character. As an example, I will train on Anna Karenina. **This model will be able to generate new text based on the text from the book!**

This network is based off of Andrej Karpathy's [post on RNNs](#) and [implementation in Torch](#). Below is the general architecture of the character-wise RNN.

First let's load in our required resources for data loading and model creation.

```
In [9]: import numpy as np
import torch
from torch import nn
import torch.nn.functional as F
```

### 1.1 Load in Data

Then, we'll load the Anna Karenina text file and convert it into integers for our network to use.

```
In [10]: # open text file and read in data as `text`
with open('data/anna.txt', 'r') as f:
    text = f.read()
```

Let's check out the first 100 characters, make sure everything is peachy. According to the [American Book Review](#), this is the 6th best first line of a book ever.

```
In [11]: text[:100]
```

```
Out[11]: 'Chapter 1\n\n\nHappy families are all alike; every unhappy family is unhappy in its ow
```

#### 1.1.1 Tokenization

In the cells, below, I'm creating a couple **dictionaries** to convert the characters to and from integers. Encoding the characters as integers makes it easier to use as input in the network.

```
In [12]: # encode the text and map each character to an integer and vice versa

# we create two dictionaries:
# 1. int2char, which maps integers to characters
```

```

# 2. char2int, which maps characters to unique integers
chars = tuple(set(text))
int2char = dict(enumerate(chars))
char2int = {ch: ii for ii, ch in int2char.items()}

# encode the text
encoded = np.array([char2int[ch] for ch in text])

```

And we can see those same characters from above, encoded as integers.

```
In [13]: encoded[:100]
```

```
Out[13]: array([26, 65,  9, 10, 47, 68,  5, 39, 72, 55, 55, 55, 30,  9, 10, 10, 35,
                39, 42,  9, 11, 66, 31, 66, 68, 53, 39,  9,  5, 68, 39,  9, 31, 31,
                39,  9, 31, 66,  0, 68, 67, 39, 68, 70, 68,  5, 35, 39, 22, 58, 65,
                 9, 10, 10, 35, 39, 42,  9, 11, 66, 31, 35, 39, 66, 53, 39, 22, 58,
                65,  9, 10, 10, 35, 39, 66, 58, 39, 66, 47, 53, 39, 32, 27, 58, 55,
                27,  9, 35, 49, 55, 55, 20, 70, 68,  5, 35, 47, 65, 66, 58])
```

## 1.2 Pre-processing the data

As you can see in our char-RNN image above, our LSTM expects an input that is **one-hot encoded** meaning that each character is converted into an integer (via our created dictionary) and *then* converted into a column vector where only it's corresponding integer index will have the value of 1 and the rest of the vector will be filled with 0's. Since we're one-hot encoding the data, let's make a function to do that!

```
In [14]: def one_hot_encode(arr, n_labels):
```

```

    # Initialize the the encoded array
    one_hot = np.zeros((np.multiply(*arr.shape), n_labels), dtype=np.float32)

    # Fill the appropriate elements with ones
    one_hot[np.arange(one_hot.shape[0]), arr.flatten()] = 1.

    # Finally reshape it to get back to the original array
    one_hot = one_hot.reshape((*arr.shape, n_labels))

    return one_hot

```

```
In [15]: # check that the function works as expected
```

```
test_seq = np.array([[3, 5, 1]])
one_hot = one_hot_encode(test_seq, 8)
```

```
print(one_hot)
```

```
[[[ 0.  0.  0.  1.  0.  0.  0.  0.]
   [ 0.  0.  0.  0.  0.  1.  0.  0.]
   [ 0.  1.  0.  0.  0.  0.  0.  0.]]]
```

## 1.3 Making training mini-batches

To train on this data, we also want to create mini-batches for training. Remember that we want our batches to be multiple sequences of some desired number of sequence steps. Considering a simple example, our batches would look like this:

In this example, we'll take the encoded characters (passed in as the `arr` parameter) and split them into multiple sequences, given by `batch_size`. Each of our sequences will be `seq_length` long.

### 1.3.1 Creating Batches

**1. The first thing we need to do is discard some of the text so we only have completely full mini-batches.**

Each batch contains  $N \times M$  characters, where  $N$  is the batch size (the number of sequences in a batch) and  $M$  is the `seq_length` or number of time steps in a sequence. Then, to get the total number of batches,  $K$ , that we can make from the array `arr`, you divide the length of `arr` by the number of characters per batch. Once you know the number of batches, you can get the total number of characters to keep from `arr`,  $N * M * K$ .

**2. After that, we need to split `arr` into  $N$  batches.**

You can do this using `arr.reshape(size)` where `size` is a tuple containing the dimensions sizes of the reshaped array. We know we want  $N$  sequences in a batch, so let's make that the size of the first dimension. For the second dimension, you can use `-1` as a placeholder in the size, it'll fill up the array with the appropriate data for you. After this, you should have an array that is  $N \times (M * K)$ .

**3. Now that we have this array, we can iterate through it to get our mini-batches.**

The idea is each batch is a  $N \times M$  window on the  $N \times (M * K)$  array. For each subsequent batch, the window moves over by `seq_length`. We also want to create both the input and target arrays. Remember that the targets are just the inputs shifted over by one character. The way I like to do this window is use `range` to take steps of size `n_steps` from 0 to `arr.shape[1]`, the total number of tokens in each sequence. That way, the integers you get from `range` always point to the start of a batch, and each window is `seq_length` wide.

**TODO:** Write the code for creating batches in the function below. The exercises in this notebook *will not be easy*. I've provided a notebook with solutions alongside this notebook. If you get stuck, checkout the solutions. The most important thing is that you don't copy and paste the code into here, **type out the solution code yourself**.

```
In [16]: def get_batches(arr, batch_size, seq_length):
        '''Create a generator that returns batches of size
            batch_size x seq_length from arr.

            Arguments
            -----
            arr: Array you want to make batches from
            batch_size: Batch size, the number of sequences per batch
            seq_length: Number of encoded chars in a sequence
        '''

        ## TODO: Get the number of batches we can make
```

```

batch_size_total = batch_size * seq_length
# total number of batches we can make
n_batches = len(arr)//batch_size_total

# Keep only enough characters to make full batches
arr = arr[:n_batches * batch_size_total]
# Reshape into batch_size rows
arr = arr.reshape((batch_size, -1))

## TODO: Iterate over the batches using a window of size seq_length
for n in range(0, arr.shape[1], seq_length):
    # The features
    x = arr[:, n:n+seq_length]
    # The targets, shifted by one
    y = np.zeros_like(x)
    try:
        y[:, :-1], y[:, -1] = x[:, 1:], arr[:, n+seq_length]
    except IndexError:
        y[:, :-1], y[:, -1] = x[:, 1:], arr[:, 0]
    yield x, y

```

### 1.3.2 Test Your Implementation

Now I'll make some data sets and we can check out what's going on as we batch data. Here, as an example, I'm going to use a batch size of 8 and 50 sequence steps.

```

In [17]: batches = get_batches(encoded, 8, 50)
        x, y = next(batches)

```

```

In [18]: # printing out the first 10 items in a sequence
        print('x\n', x[:10, :10])
        print('\ny\n', y[:10, :10])

```

```

x
[[26 65  9 10 47 68  5 39 72 55]
 [53 32 58 39 47 65  9 47 39  9]
 [68 58 64 39 32  5 39  9 39 42]
 [53 39 47 65 68 39 73 65 66 68]
 [39 53  9 27 39 65 68  5 39 47]
 [73 22 53 53 66 32 58 39  9 58]
 [39 56 58 58  9 39 65  9 64 39]
 [69 12 31 32 58 53  0 35 49 39]]

```

```

y
[[65  9 10 47 68  5 39 72 55 55]
 [32 58 39 47 65  9 47 39  9 47]
 [58 64 39 32  5 39  9 39 42 32]
 [39 47 65 68 39 73 65 66 68 42]
 [53  9 27 39 65 68  5 39 47 68]

```

```
[22 53 53 66 32 58 39  9 58 64]
[56 58 58  9 39 65  9 64 39 53]
[12 31 32 58 53  0 35 49 39 50]]
```

If you implemented `get_batches` correctly, the above output should look something like ““ x [[25 8 60 11 45 27 28 73 1 2]][17 7 20 73 45 8 60 45 73 60] [27 20 80 73 7 28 73 60 73 65]][17 73 45 8 27 73 66 8 46 27] [73 17 60 12 73 8 27 28 73 45]][66 64 17 17 46 7 20 73 60 20] [73 76 20 20 60 73 8 60 80 73]][47 35 43 7 20 17 24 50 37 73]]

y [[ 8 60 11 45 27 28 73 1 2 2]][ 7 20 73 45 8 60 45 73 60 45] [20 80 73 7 28 73 60 73 65 7]][73 45 8 27 73 66 8 46 27 65] [17 60 12 73 8 27 28 73 45 27]][64 17 17 46 7 20 73 60 20 80] [76 20 20 60 73 8 60 80 73 17]][35 43 7 20 17 24 50 37 73 36]] “although the exact numbers may be different. Check to make sure the data is shifted over one step for y’.

---

## 1.4 Defining the network with PyTorch

Below is where you’ll define the network.

Next, you’ll use PyTorch to define the architecture of the network. We start by defining the layers and operations we want. Then, define a method for the forward pass. You’ve also been given a method for predicting characters.

### 1.4.1 Model Structure

In `__init__` the suggested structure is as follows: \* Create and store the necessary dictionaries (this has been done for you) \* Define an LSTM layer that takes as params: an input size (the number of characters), a hidden layer size `n_hidden`, a number of layers `n_layers`, a dropout probability `drop_prob`, and a `batch_first` boolean (True, since we are batching) \* Define a dropout layer with `dropout_prob` \* Define a fully-connected layer with params: input size `n_hidden` and output size (the number of characters) \* Finally, initialize the weights (again, this has been given)

Note that some parameters have been named and given in the `__init__` function, and we use them and store them by doing something like `self.drop_prob = drop_prob`.

---

### 1.4.2 LSTM Inputs/Outputs

You can create a basic [LSTM layer](#) as follows

```
self.lstm = nn.LSTM(input_size, n_hidden, n_layers,
                    dropout=drop_prob, batch_first=True)
```

where `input_size` is the number of characters this cell expects to see as sequential input, and `n_hidden` is the number of units in the hidden layers in the cell. And we can add dropout by adding a dropout parameter with a specified probability; this will automatically add dropout to the inputs or outputs. Finally, in the forward function, we can stack up the LSTM cells into layers using `.view`. With this, you pass in a list of cells and it will send the output of one cell into the next cell.

We also need to create an initial hidden state of all zeros. This is done like so

```
self.init_hidden()
```

```
In [19]: # check if GPU is available
train_on_gpu = torch.cuda.is_available()
if(train_on_gpu):
    print('Training on GPU!')
else:
    print('No GPU available, training on CPU; consider making n_epochs very small.')
```

Training on GPU!

```
In [20]: class CharRNN(nn.Module):

    def __init__(self, tokens, n_hidden=256, n_layers=2,
                  drop_prob=0.5, lr=0.001):
        super().__init__()
        self.drop_prob = drop_prob
        self.n_layers = n_layers
        self.n_hidden = n_hidden
        self.lr = lr

        # creating character dictionaries
        self.chars = tokens
        self.int2char = dict(enumerate(self.chars))
        self.char2int = {ch: ii for ii, ch in self.int2char.items()}

        ## TODO: define the layers of the model
        self.lstm = nn.LSTM(len(self.chars), n_hidden, n_layers,
                             dropout=drop_prob, batch_first=True)
        self.dropout = nn.Dropout(drop_prob)

        ## TODO: define the final, fully-connected output layer
        self.fc = nn.Linear(n_hidden, len(self.chars))

    def forward(self, x, hidden):
        ''' Forward pass through the network.
            These inputs are x, and the hidden/cell state `hidden`. '''

        ## TODO: Get the outputs and the new hidden state from the lstm
        r_output, hidden = self.lstm(x, hidden)

        ## TODO: pass through a dropout layer
        out = self.dropout(r_output)

        # Stack up LSTM outputs using view
        # you may need to use contiguous to reshape the output
        out = out.contiguous().view(-1, self.n_hidden)
```

```

    ## TODO: put x through the fully-connected layer
    out = self.fc(out)

    # return the final output and the hidden state
    # return the final output and the hidden state
    return out, hidden

def init_hidden(self, batch_size):
    ''' Initializes hidden state '''
    # Create two new tensors with sizes n_layers x batch_size x n_hidden,
    # initialized to zero, for hidden state and cell state of LSTM
    weight = next(self.parameters()).data

    if (train_on_gpu):
        hidden = (weight.new(self.n_layers, batch_size, self.n_hidden).zero_().cuda(),
                  weight.new(self.n_layers, batch_size, self.n_hidden).zero_().cuda())
    else:
        hidden = (weight.new(self.n_layers, batch_size, self.n_hidden).zero_(),
                  weight.new(self.n_layers, batch_size, self.n_hidden).zero_())

    return hidden

```

## 1.5 Time to train

The train function gives us the ability to set the number of epochs, the learning rate, and other parameters.

Below we're using an Adam optimizer and cross entropy loss since we are looking at character class scores as output. We calculate the loss and perform backpropagation, as usual!

A couple of details about training: >\* Within the batch loop, we detach the hidden state from its history; this time setting it equal to a new *tuple* variable because an LSTM has a hidden state that is a tuple of the hidden and cell states. \* We use `clip_grad_norm` to help prevent exploding gradients.

```

In [21]: def train(net, data, epochs=10, batch_size=10, seq_length=50, lr=0.001, clip=5, val_fraction=0.1):
    ''' Training a network

    Arguments
    -----

    net: CharRNN network
    data: text data to train the network
    epochs: Number of epochs to train
    batch_size: Number of mini-sequences per mini-batch, aka batch size
    seq_length: Number of character steps per mini-batch
    lr: learning rate

```

```

        clip: gradient clipping
        val_frac: Fraction of data to hold out for validation
        print_every: Number of steps for printing training and validation loss

'''
net.train()

opt = torch.optim.Adam(net.parameters(), lr=lr)
criterion = nn.CrossEntropyLoss()

# create training and validation data
val_idx = int(len(data)*(1-val_frac))
data, val_data = data[:val_idx], data[val_idx:]

if(train_on_gpu):
    net.cuda()

counter = 0
n_chars = len(net.chars)
for e in range(epochs):
    # initialize hidden state
    h = net.init_hidden(batch_size)

    for x, y in get_batches(data, batch_size, seq_length):
        counter += 1

        # One-hot encode our data and make them Torch tensors
        x = one_hot_encode(x, n_chars)
        inputs, targets = torch.from_numpy(x), torch.from_numpy(y)

        if(train_on_gpu):
            inputs, targets = inputs.cuda(), targets.cuda()

        # Creating new variables for the hidden state, otherwise
        # we'd backprop through the entire training history
        h = tuple([each.data for each in h])

        # zero accumulated gradients
        net.zero_grad()

        # get the output from the model
        output, h = net(inputs, h)

        # calculate the loss and perform backprop
        loss = criterion(output, targets.view(batch_size*seq_length))
        loss.backward()
        # `clip_grad_norm` helps prevent the exploding gradient problem in RNNs / L
        nn.utils.clip_grad_norm_(net.parameters(), clip)

```



```

opt.step()

# loss stats
if counter % print_every == 0:
    # Get validation loss
    val_h = net.init_hidden(batch_size)
    val_losses = []
    net.eval()
    for x, y in get_batches(val_data, batch_size, seq_length):
        # One-hot encode our data and make them Torch tensors
        x = one_hot_encode(x, n_chars)
        x, y = torch.from_numpy(x), torch.from_numpy(y)

        # Creating new variables for the hidden state, otherwise
        # we'd backprop through the entire training history
        val_h = tuple([each.data for each in val_h])

        inputs, targets = x, y
        if(train_on_gpu):
            inputs, targets = inputs.cuda(), targets.cuda()

        output, val_h = net(inputs, val_h)
        val_loss = criterion(output, targets.view(batch_size*seq_length))

        val_losses.append(val_loss.item())

    net.train() # reset to train mode after iterationg through validation d

    print("Epoch: {}/{}...".format(e+1, epochs),
          "Step: {}...".format(counter),
          "Loss: {:.4f}...".format(loss.item()),
          "Val Loss: {:.4f}".format(np.mean(val_losses)))

```

## 1.6 Instantiating the model

Now we can actually train the network. First we'll create the network itself, with some given hyperparameters. Then, define the mini-batches sizes, and start training!

```

In [24]: ## TODO: set you model hyperparameters
         # define and print the net
         n_hidden= 512
         n_layers= 2

         net = CharRNN(chars, n_hidden, n_layers)
         print(net)

```

```

CharRNN(
  (lstm): LSTM(83, 512, num_layers=2, batch_first=True, dropout=0.5)
)

```

```

(dropout): Dropout(p=0.5)
(fc): Linear(in_features=512, out_features=83, bias=True)
)

```

### 1.6.1 Set your training hyperparameters!

```

In [26]: batch_size = 128
        seq_length = 100
        n_epochs = 20 # start small if you are just testing initial behavior

        # train the model
        train(net, encoded, epochs=n_epochs, batch_size=batch_size, seq_length=seq_length, lr=0

Epoch: 1/20... Step: 10... Loss: 3.2522... Val Loss: 3.1988
Epoch: 1/20... Step: 20... Loss: 3.1508... Val Loss: 3.1361
Epoch: 1/20... Step: 30... Loss: 3.1385... Val Loss: 3.1237
Epoch: 1/20... Step: 40... Loss: 3.1096... Val Loss: 3.1188
Epoch: 1/20... Step: 50... Loss: 3.1429... Val Loss: 3.1164
Epoch: 1/20... Step: 60... Loss: 3.1152... Val Loss: 3.1128
Epoch: 1/20... Step: 70... Loss: 3.1037... Val Loss: 3.1073
Epoch: 1/20... Step: 80... Loss: 3.1056... Val Loss: 3.0927
Epoch: 1/20... Step: 90... Loss: 3.0856... Val Loss: 3.0627
Epoch: 1/20... Step: 100... Loss: 3.0129... Val Loss: 3.0004
Epoch: 1/20... Step: 110... Loss: 2.9536... Val Loss: 2.9232
Epoch: 1/20... Step: 120... Loss: 2.8469... Val Loss: 2.8184
Epoch: 1/20... Step: 130... Loss: 2.7618... Val Loss: 2.7188
Epoch: 2/20... Step: 140... Loss: 2.6744... Val Loss: 2.6147
Epoch: 2/20... Step: 150... Loss: 2.5986... Val Loss: 2.5390
Epoch: 2/20... Step: 160... Loss: 2.5237... Val Loss: 2.4842
Epoch: 2/20... Step: 170... Loss: 2.4572... Val Loss: 2.4421
Epoch: 2/20... Step: 180... Loss: 2.4335... Val Loss: 2.4115
Epoch: 2/20... Step: 190... Loss: 2.3801... Val Loss: 2.3840
Epoch: 2/20... Step: 200... Loss: 2.3820... Val Loss: 2.3502
Epoch: 2/20... Step: 210... Loss: 2.3461... Val Loss: 2.3246
Epoch: 2/20... Step: 220... Loss: 2.3090... Val Loss: 2.2977
Epoch: 2/20... Step: 230... Loss: 2.3087... Val Loss: 2.2684
Epoch: 2/20... Step: 240... Loss: 2.2781... Val Loss: 2.2419
Epoch: 2/20... Step: 250... Loss: 2.2075... Val Loss: 2.2138
Epoch: 2/20... Step: 260... Loss: 2.1822... Val Loss: 2.1889
Epoch: 2/20... Step: 270... Loss: 2.2052... Val Loss: 2.1652
Epoch: 3/20... Step: 280... Loss: 2.1913... Val Loss: 2.1495
Epoch: 3/20... Step: 290... Loss: 2.1486... Val Loss: 2.1239
Epoch: 3/20... Step: 300... Loss: 2.1233... Val Loss: 2.1008
Epoch: 3/20... Step: 310... Loss: 2.1052... Val Loss: 2.0816
Epoch: 3/20... Step: 320... Loss: 2.0794... Val Loss: 2.0625
Epoch: 3/20... Step: 330... Loss: 2.0409... Val Loss: 2.0468
Epoch: 3/20... Step: 340... Loss: 2.0721... Val Loss: 2.0280

```

Epoch: 3/20... Step: 350... Loss: 2.0434... Val Loss: 2.0115  
 Epoch: 3/20... Step: 360... Loss: 1.9774... Val Loss: 1.9923  
 Epoch: 3/20... Step: 370... Loss: 2.0063... Val Loss: 1.9770  
 Epoch: 3/20... Step: 380... Loss: 1.9851... Val Loss: 1.9636  
 Epoch: 3/20... Step: 390... Loss: 1.9554... Val Loss: 1.9478  
 Epoch: 3/20... Step: 400... Loss: 1.9162... Val Loss: 1.9297  
 Epoch: 3/20... Step: 410... Loss: 1.9438... Val Loss: 1.9146  
 Epoch: 4/20... Step: 420... Loss: 1.9319... Val Loss: 1.8999  
 Epoch: 4/20... Step: 430... Loss: 1.9168... Val Loss: 1.8827  
 Epoch: 4/20... Step: 440... Loss: 1.8993... Val Loss: 1.8764  
 Epoch: 4/20... Step: 450... Loss: 1.8465... Val Loss: 1.8616  
 Epoch: 4/20... Step: 460... Loss: 1.8314... Val Loss: 1.8504  
 Epoch: 4/20... Step: 470... Loss: 1.8680... Val Loss: 1.8375  
 Epoch: 4/20... Step: 480... Loss: 1.8426... Val Loss: 1.8257  
 Epoch: 4/20... Step: 490... Loss: 1.8480... Val Loss: 1.8128  
 Epoch: 4/20... Step: 500... Loss: 1.8502... Val Loss: 1.8080  
 Epoch: 4/20... Step: 510... Loss: 1.8173... Val Loss: 1.7950  
 Epoch: 4/20... Step: 520... Loss: 1.8304... Val Loss: 1.7874  
 Epoch: 4/20... Step: 530... Loss: 1.7867... Val Loss: 1.7756  
 Epoch: 4/20... Step: 540... Loss: 1.7472... Val Loss: 1.7647  
 Epoch: 4/20... Step: 550... Loss: 1.7942... Val Loss: 1.7511  
 Epoch: 5/20... Step: 560... Loss: 1.7632... Val Loss: 1.7466  
 Epoch: 5/20... Step: 570... Loss: 1.7513... Val Loss: 1.7309  
 Epoch: 5/20... Step: 580... Loss: 1.7309... Val Loss: 1.7239  
 Epoch: 5/20... Step: 590... Loss: 1.7316... Val Loss: 1.7137  
 Epoch: 5/20... Step: 600... Loss: 1.7219... Val Loss: 1.7129  
 Epoch: 5/20... Step: 610... Loss: 1.7018... Val Loss: 1.7007  
 Epoch: 5/20... Step: 620... Loss: 1.7054... Val Loss: 1.6934  
 Epoch: 5/20... Step: 630... Loss: 1.7218... Val Loss: 1.6881  
 Epoch: 5/20... Step: 640... Loss: 1.6802... Val Loss: 1.6809  
 Epoch: 5/20... Step: 650... Loss: 1.6816... Val Loss: 1.6696  
 Epoch: 5/20... Step: 660... Loss: 1.6584... Val Loss: 1.6671  
 Epoch: 5/20... Step: 670... Loss: 1.6823... Val Loss: 1.6567  
 Epoch: 5/20... Step: 680... Loss: 1.6795... Val Loss: 1.6537  
 Epoch: 5/20... Step: 690... Loss: 1.6573... Val Loss: 1.6485  
 Epoch: 6/20... Step: 700... Loss: 1.6501... Val Loss: 1.6481  
 Epoch: 6/20... Step: 710... Loss: 1.6389... Val Loss: 1.6321  
 Epoch: 6/20... Step: 720... Loss: 1.6251... Val Loss: 1.6286  
 Epoch: 6/20... Step: 730... Loss: 1.6528... Val Loss: 1.6222  
 Epoch: 6/20... Step: 740... Loss: 1.6135... Val Loss: 1.6186  
 Epoch: 6/20... Step: 750... Loss: 1.5928... Val Loss: 1.6157  
 Epoch: 6/20... Step: 760... Loss: 1.6316... Val Loss: 1.6075  
 Epoch: 6/20... Step: 770... Loss: 1.6082... Val Loss: 1.6043  
 Epoch: 6/20... Step: 780... Loss: 1.5921... Val Loss: 1.5952  
 Epoch: 6/20... Step: 790... Loss: 1.5852... Val Loss: 1.5907  
 Epoch: 6/20... Step: 800... Loss: 1.5961... Val Loss: 1.5875  
 Epoch: 6/20... Step: 810... Loss: 1.5901... Val Loss: 1.5815  
 Epoch: 6/20... Step: 820... Loss: 1.5511... Val Loss: 1.5781

Epoch: 6/20... Step: 830... Loss: 1.5911... Val Loss: 1.5727  
Epoch: 7/20... Step: 840... Loss: 1.5576... Val Loss: 1.5709  
Epoch: 7/20... Step: 850... Loss: 1.5641... Val Loss: 1.5623  
Epoch: 7/20... Step: 860... Loss: 1.5481... Val Loss: 1.5579  
Epoch: 7/20... Step: 870... Loss: 1.5580... Val Loss: 1.5521  
Epoch: 7/20... Step: 880... Loss: 1.5562... Val Loss: 1.5486  
Epoch: 7/20... Step: 890... Loss: 1.5523... Val Loss: 1.5447  
Epoch: 7/20... Step: 900... Loss: 1.5352... Val Loss: 1.5442  
Epoch: 7/20... Step: 910... Loss: 1.5031... Val Loss: 1.5392  
Epoch: 7/20... Step: 920... Loss: 1.5229... Val Loss: 1.5347  
Epoch: 7/20... Step: 930... Loss: 1.5180... Val Loss: 1.5314  
Epoch: 7/20... Step: 940... Loss: 1.5288... Val Loss: 1.5298  
Epoch: 7/20... Step: 950... Loss: 1.5325... Val Loss: 1.5231  
Epoch: 7/20... Step: 960... Loss: 1.5330... Val Loss: 1.5172  
Epoch: 7/20... Step: 970... Loss: 1.5455... Val Loss: 1.5161  
Epoch: 8/20... Step: 980... Loss: 1.5061... Val Loss: 1.5160  
Epoch: 8/20... Step: 990... Loss: 1.5090... Val Loss: 1.5094  
Epoch: 8/20... Step: 1000... Loss: 1.5036... Val Loss: 1.5050  
Epoch: 8/20... Step: 1010... Loss: 1.5457... Val Loss: 1.5019  
Epoch: 8/20... Step: 1020... Loss: 1.4956... Val Loss: 1.5000  
Epoch: 8/20... Step: 1030... Loss: 1.4934... Val Loss: 1.4942  
Epoch: 8/20... Step: 1040... Loss: 1.5029... Val Loss: 1.4967  
Epoch: 8/20... Step: 1050... Loss: 1.4795... Val Loss: 1.4919  
Epoch: 8/20... Step: 1060... Loss: 1.4825... Val Loss: 1.4848  
Epoch: 8/20... Step: 1070... Loss: 1.4863... Val Loss: 1.4873  
Epoch: 8/20... Step: 1080... Loss: 1.4870... Val Loss: 1.4811  
Epoch: 8/20... Step: 1090... Loss: 1.4696... Val Loss: 1.4801  
Epoch: 8/20... Step: 1100... Loss: 1.4580... Val Loss: 1.4782  
Epoch: 8/20... Step: 1110... Loss: 1.4699... Val Loss: 1.4715  
Epoch: 9/20... Step: 1120... Loss: 1.4784... Val Loss: 1.4717  
Epoch: 9/20... Step: 1130... Loss: 1.4728... Val Loss: 1.4701  
Epoch: 9/20... Step: 1140... Loss: 1.4786... Val Loss: 1.4640  
Epoch: 9/20... Step: 1150... Loss: 1.4826... Val Loss: 1.4629  
Epoch: 9/20... Step: 1160... Loss: 1.4484... Val Loss: 1.4603  
Epoch: 9/20... Step: 1170... Loss: 1.4532... Val Loss: 1.4588  
Epoch: 9/20... Step: 1180... Loss: 1.4551... Val Loss: 1.4585  
Epoch: 9/20... Step: 1190... Loss: 1.4827... Val Loss: 1.4567  
Epoch: 9/20... Step: 1200... Loss: 1.4282... Val Loss: 1.4544  
Epoch: 9/20... Step: 1210... Loss: 1.4341... Val Loss: 1.4494  
Epoch: 9/20... Step: 1220... Loss: 1.4433... Val Loss: 1.4485  
Epoch: 9/20... Step: 1230... Loss: 1.4140... Val Loss: 1.4446  
Epoch: 9/20... Step: 1240... Loss: 1.4335... Val Loss: 1.4396  
Epoch: 9/20... Step: 1250... Loss: 1.4350... Val Loss: 1.4393  
Epoch: 10/20... Step: 1260... Loss: 1.4476... Val Loss: 1.4385  
Epoch: 10/20... Step: 1270... Loss: 1.4319... Val Loss: 1.4372  
Epoch: 10/20... Step: 1280... Loss: 1.4371... Val Loss: 1.4335  
Epoch: 10/20... Step: 1290... Loss: 1.4334... Val Loss: 1.4296  
Epoch: 10/20... Step: 1300... Loss: 1.4257... Val Loss: 1.4303

Epoch: 10/20... Step: 1310... Loss: 1.4206... Val Loss: 1.4285  
 Epoch: 10/20... Step: 1320... Loss: 1.4018... Val Loss: 1.4247  
 Epoch: 10/20... Step: 1330... Loss: 1.4100... Val Loss: 1.4263  
 Epoch: 10/20... Step: 1340... Loss: 1.3882... Val Loss: 1.4244  
 Epoch: 10/20... Step: 1350... Loss: 1.3872... Val Loss: 1.4202  
 Epoch: 10/20... Step: 1360... Loss: 1.3883... Val Loss: 1.4197  
 Epoch: 10/20... Step: 1370... Loss: 1.3826... Val Loss: 1.4151  
 Epoch: 10/20... Step: 1380... Loss: 1.4113... Val Loss: 1.4158  
 Epoch: 10/20... Step: 1390... Loss: 1.4241... Val Loss: 1.4141  
 Epoch: 11/20... Step: 1400... Loss: 1.4201... Val Loss: 1.4115  
 Epoch: 11/20... Step: 1410... Loss: 1.4378... Val Loss: 1.4098  
 Epoch: 11/20... Step: 1420... Loss: 1.4223... Val Loss: 1.4059  
 Epoch: 11/20... Step: 1430... Loss: 1.3954... Val Loss: 1.4061  
 Epoch: 11/20... Step: 1440... Loss: 1.4196... Val Loss: 1.4092  
 Epoch: 11/20... Step: 1450... Loss: 1.3541... Val Loss: 1.4041  
 Epoch: 11/20... Step: 1460... Loss: 1.3701... Val Loss: 1.4017  
 Epoch: 11/20... Step: 1470... Loss: 1.3597... Val Loss: 1.4046  
 Epoch: 11/20... Step: 1480... Loss: 1.3903... Val Loss: 1.3985  
 Epoch: 11/20... Step: 1490... Loss: 1.3702... Val Loss: 1.3998  
 Epoch: 11/20... Step: 1500... Loss: 1.3523... Val Loss: 1.4003  
 Epoch: 11/20... Step: 1510... Loss: 1.3410... Val Loss: 1.3969  
 Epoch: 11/20... Step: 1520... Loss: 1.3842... Val Loss: 1.3902  
 Epoch: 12/20... Step: 1530... Loss: 1.4396... Val Loss: 1.3932  
 Epoch: 12/20... Step: 1540... Loss: 1.3905... Val Loss: 1.3918  
 Epoch: 12/20... Step: 1550... Loss: 1.3862... Val Loss: 1.3886  
 Epoch: 12/20... Step: 1560... Loss: 1.4039... Val Loss: 1.3848  
 Epoch: 12/20... Step: 1570... Loss: 1.3510... Val Loss: 1.3891  
 Epoch: 12/20... Step: 1580... Loss: 1.3249... Val Loss: 1.3870  
 Epoch: 12/20... Step: 1590... Loss: 1.3292... Val Loss: 1.3857  
 Epoch: 12/20... Step: 1600... Loss: 1.3531... Val Loss: 1.3853  
 Epoch: 12/20... Step: 1610... Loss: 1.3385... Val Loss: 1.3851  
 Epoch: 12/20... Step: 1620... Loss: 1.3482... Val Loss: 1.3783  
 Epoch: 12/20... Step: 1630... Loss: 1.3660... Val Loss: 1.3799  
 Epoch: 12/20... Step: 1640... Loss: 1.3424... Val Loss: 1.3796  
 Epoch: 12/20... Step: 1650... Loss: 1.3226... Val Loss: 1.3772  
 Epoch: 12/20... Step: 1660... Loss: 1.3715... Val Loss: 1.3730  
 Epoch: 13/20... Step: 1670... Loss: 1.3464... Val Loss: 1.3769  
 Epoch: 13/20... Step: 1680... Loss: 1.3571... Val Loss: 1.3725  
 Epoch: 13/20... Step: 1690... Loss: 1.3290... Val Loss: 1.3684  
 Epoch: 13/20... Step: 1700... Loss: 1.3394... Val Loss: 1.3681  
 Epoch: 13/20... Step: 1710... Loss: 1.3091... Val Loss: 1.3700  
 Epoch: 13/20... Step: 1720... Loss: 1.3207... Val Loss: 1.3725  
 Epoch: 13/20... Step: 1730... Loss: 1.3595... Val Loss: 1.3679  
 Epoch: 13/20... Step: 1740... Loss: 1.3308... Val Loss: 1.3696  
 Epoch: 13/20... Step: 1750... Loss: 1.2979... Val Loss: 1.3715  
 Epoch: 13/20... Step: 1760... Loss: 1.3292... Val Loss: 1.3654  
 Epoch: 13/20... Step: 1770... Loss: 1.3431... Val Loss: 1.3649  
 Epoch: 13/20... Step: 1780... Loss: 1.3191... Val Loss: 1.3620

Epoch: 13/20... Step: 1790... Loss: 1.3125... Val Loss: 1.3586  
 Epoch: 13/20... Step: 1800... Loss: 1.3356... Val Loss: 1.3573  
 Epoch: 14/20... Step: 1810... Loss: 1.3378... Val Loss: 1.3682  
 Epoch: 14/20... Step: 1820... Loss: 1.3216... Val Loss: 1.3608  
 Epoch: 14/20... Step: 1830... Loss: 1.3362... Val Loss: 1.3560  
 Epoch: 14/20... Step: 1840... Loss: 1.2818... Val Loss: 1.3564  
 Epoch: 14/20... Step: 1850... Loss: 1.2700... Val Loss: 1.3583  
 Epoch: 14/20... Step: 1860... Loss: 1.3295... Val Loss: 1.3581  
 Epoch: 14/20... Step: 1870... Loss: 1.3361... Val Loss: 1.3550  
 Epoch: 14/20... Step: 1880... Loss: 1.3298... Val Loss: 1.3570  
 Epoch: 14/20... Step: 1890... Loss: 1.3335... Val Loss: 1.3592  
 Epoch: 14/20... Step: 1900... Loss: 1.3168... Val Loss: 1.3528  
 Epoch: 14/20... Step: 1910... Loss: 1.3158... Val Loss: 1.3505  
 Epoch: 14/20... Step: 1920... Loss: 1.3114... Val Loss: 1.3487  
 Epoch: 14/20... Step: 1930... Loss: 1.2777... Val Loss: 1.3482  
 Epoch: 14/20... Step: 1940... Loss: 1.3411... Val Loss: 1.3455  
 Epoch: 15/20... Step: 1950... Loss: 1.3011... Val Loss: 1.3526  
 Epoch: 15/20... Step: 1960... Loss: 1.3114... Val Loss: 1.3483  
 Epoch: 15/20... Step: 1970... Loss: 1.3029... Val Loss: 1.3448  
 Epoch: 15/20... Step: 1980... Loss: 1.3048... Val Loss: 1.3474  
 Epoch: 15/20... Step: 1990... Loss: 1.2911... Val Loss: 1.3461  
 Epoch: 15/20... Step: 2000... Loss: 1.2755... Val Loss: 1.3434  
 Epoch: 15/20... Step: 2010... Loss: 1.2984... Val Loss: 1.3419  
 Epoch: 15/20... Step: 2020... Loss: 1.3062... Val Loss: 1.3449  
 Epoch: 15/20... Step: 2030... Loss: 1.2871... Val Loss: 1.3454  
 Epoch: 15/20... Step: 2040... Loss: 1.2962... Val Loss: 1.3373  
 Epoch: 15/20... Step: 2050... Loss: 1.2870... Val Loss: 1.3399  
 Epoch: 15/20... Step: 2060... Loss: 1.3063... Val Loss: 1.3383  
 Epoch: 15/20... Step: 2070... Loss: 1.2991... Val Loss: 1.3341  
 Epoch: 15/20... Step: 2080... Loss: 1.2937... Val Loss: 1.3327  
 Epoch: 16/20... Step: 2090... Loss: 1.3005... Val Loss: 1.3354  
 Epoch: 16/20... Step: 2100... Loss: 1.2800... Val Loss: 1.3318  
 Epoch: 16/20... Step: 2110... Loss: 1.2695... Val Loss: 1.3323  
 Epoch: 16/20... Step: 2120... Loss: 1.2888... Val Loss: 1.3313  
 Epoch: 16/20... Step: 2130... Loss: 1.2711... Val Loss: 1.3325  
 Epoch: 16/20... Step: 2140... Loss: 1.2755... Val Loss: 1.3310  
 Epoch: 16/20... Step: 2150... Loss: 1.3072... Val Loss: 1.3322  
 Epoch: 16/20... Step: 2160... Loss: 1.2828... Val Loss: 1.3352  
 Epoch: 16/20... Step: 2170... Loss: 1.2788... Val Loss: 1.3330  
 Epoch: 16/20... Step: 2180... Loss: 1.2791... Val Loss: 1.3289  
 Epoch: 16/20... Step: 2190... Loss: 1.2980... Val Loss: 1.3279  
 Epoch: 16/20... Step: 2200... Loss: 1.2716... Val Loss: 1.3303  
 Epoch: 16/20... Step: 2210... Loss: 1.2362... Val Loss: 1.3253  
 Epoch: 16/20... Step: 2220... Loss: 1.2877... Val Loss: 1.3220  
 Epoch: 17/20... Step: 2230... Loss: 1.2547... Val Loss: 1.3254  
 Epoch: 17/20... Step: 2240... Loss: 1.2697... Val Loss: 1.3258  
 Epoch: 17/20... Step: 2250... Loss: 1.2535... Val Loss: 1.3261  
 Epoch: 17/20... Step: 2260... Loss: 1.2569... Val Loss: 1.3278

Epoch: 17/20... Step: 2270... Loss: 1.2713... Val Loss: 1.3239  
Epoch: 17/20... Step: 2280... Loss: 1.2852... Val Loss: 1.3212  
Epoch: 17/20... Step: 2290... Loss: 1.2796... Val Loss: 1.3215  
Epoch: 17/20... Step: 2300... Loss: 1.2470... Val Loss: 1.3250  
Epoch: 17/20... Step: 2310... Loss: 1.2703... Val Loss: 1.3214  
Epoch: 17/20... Step: 2320... Loss: 1.2638... Val Loss: 1.3193  
Epoch: 17/20... Step: 2330... Loss: 1.2474... Val Loss: 1.3173  
Epoch: 17/20... Step: 2340... Loss: 1.2733... Val Loss: 1.3177  
Epoch: 17/20... Step: 2350... Loss: 1.2734... Val Loss: 1.3151  
Epoch: 17/20... Step: 2360... Loss: 1.2726... Val Loss: 1.3150  
Epoch: 18/20... Step: 2370... Loss: 1.2498... Val Loss: 1.3163  
Epoch: 18/20... Step: 2380... Loss: 1.2560... Val Loss: 1.3183  
Epoch: 18/20... Step: 2390... Loss: 1.2576... Val Loss: 1.3183  
Epoch: 18/20... Step: 2400... Loss: 1.2845... Val Loss: 1.3179  
Epoch: 18/20... Step: 2410... Loss: 1.2796... Val Loss: 1.3152  
Epoch: 18/20... Step: 2420... Loss: 1.2615... Val Loss: 1.3152  
Epoch: 18/20... Step: 2430... Loss: 1.2640... Val Loss: 1.3158  
Epoch: 18/20... Step: 2440... Loss: 1.2495... Val Loss: 1.3169  
Epoch: 18/20... Step: 2450... Loss: 1.2379... Val Loss: 1.3135  
Epoch: 18/20... Step: 2460... Loss: 1.2698... Val Loss: 1.3128  
Epoch: 18/20... Step: 2470... Loss: 1.2439... Val Loss: 1.3067  
Epoch: 18/20... Step: 2480... Loss: 1.2540... Val Loss: 1.3106  
Epoch: 18/20... Step: 2490... Loss: 1.2418... Val Loss: 1.3057  
Epoch: 18/20... Step: 2500... Loss: 1.2449... Val Loss: 1.3069  
Epoch: 19/20... Step: 2510... Loss: 1.2530... Val Loss: 1.3089  
Epoch: 19/20... Step: 2520... Loss: 1.2552... Val Loss: 1.3126  
Epoch: 19/20... Step: 2530... Loss: 1.2663... Val Loss: 1.3100  
Epoch: 19/20... Step: 2540... Loss: 1.2717... Val Loss: 1.3105  
Epoch: 19/20... Step: 2550... Loss: 1.2419... Val Loss: 1.3091  
Epoch: 19/20... Step: 2560... Loss: 1.2529... Val Loss: 1.3057  
Epoch: 19/20... Step: 2570... Loss: 1.2497... Val Loss: 1.3048  
Epoch: 19/20... Step: 2580... Loss: 1.2710... Val Loss: 1.3086  
Epoch: 19/20... Step: 2590... Loss: 1.2252... Val Loss: 1.3055  
Epoch: 19/20... Step: 2600... Loss: 1.2245... Val Loss: 1.3028  
Epoch: 19/20... Step: 2610... Loss: 1.2436... Val Loss: 1.3053  
Epoch: 19/20... Step: 2620... Loss: 1.2266... Val Loss: 1.3076  
Epoch: 19/20... Step: 2630... Loss: 1.2293... Val Loss: 1.2981  
Epoch: 19/20... Step: 2640... Loss: 1.2467... Val Loss: 1.2956  
Epoch: 20/20... Step: 2650... Loss: 1.2556... Val Loss: 1.3012  
Epoch: 20/20... Step: 2660... Loss: 1.2496... Val Loss: 1.3075  
Epoch: 20/20... Step: 2670... Loss: 1.2435... Val Loss: 1.3042  
Epoch: 20/20... Step: 2680... Loss: 1.2509... Val Loss: 1.3043  
Epoch: 20/20... Step: 2690... Loss: 1.2397... Val Loss: 1.3036  
Epoch: 20/20... Step: 2700... Loss: 1.2476... Val Loss: 1.2999  
Epoch: 20/20... Step: 2710... Loss: 1.2180... Val Loss: 1.3019  
Epoch: 20/20... Step: 2720... Loss: 1.2103... Val Loss: 1.3062  
Epoch: 20/20... Step: 2730... Loss: 1.2150... Val Loss: 1.3017  
Epoch: 20/20... Step: 2740... Loss: 1.2109... Val Loss: 1.3009